




August 8, 2026

ProofyX

Is this real?

ProofyX answers that in seconds — for any image, video or voice — and shows the evidence behind the answer.



Investor Pitch Deck

Seed Round

Nobody can tell what is real any more.

Synthetic media crossed the threshold where human review stops working — and the losses are already being booked.

8M+

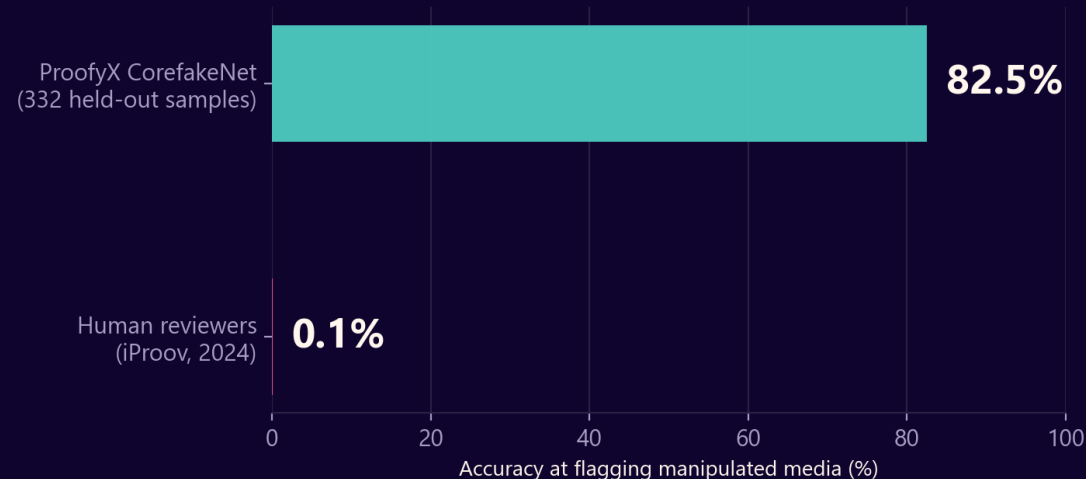
Deepfake files in circulation in 2025,
up from 500K in 2023

\$3B+

Lost to deepfake-enabled fraud in the US,
January to September 2025

\$280K

Average loss per single
deepfake fraud incident



Detection is now a compliance problem, not a curiosity. Banks, newsrooms and courts need an answer they can put in a file.



Detection just became **legally mandatory.**

Regulation is converting deepfake detection from a nice-to-have into compulsory spend — on a clock we can name.

EU AI ACT

Article 50

From 2 August 2026, AI-generated content must be machine-readable and detectable. Article 99 penalties reach EUR 15M or 3% of global turnover.

UNITED STATES

TAKE IT DOWN Act + 46 states

Signed May 2025. 169 state deepfake laws since 2022, 146 bills in 2025 alone. FinCEN warned financial institutions in Q4 2025.

INDIA

IT Rules amendment, 2026

Synthetic-media labelling and takedown duties for intermediaries — the market where our first design partners sit.

GLOBAL

China, South Korea, FATF

FATF named deepfakes a direct threat to AML and customer due diligence. South Korea legislated after a 1,625% spike in deepfake fraud.

The buyer no longer has to be convinced the problem is real — a regulator already told them. That changes the sales cycle from evangelism to procurement.

One number, and the reason behind it.

ProofyX returns a 0-100 Trust Score for any image, video or audio file — plus the heatmap and per-model evidence that produced it.

EVIDENCE, NOT A VERDICT

Explainable by construction

GradCAM heatmaps show which pixels drove the score, and every model in the ensemble reports separately. Competitors return a binary label; a compliance officer cannot file a binary label.

BREADTH

Image, video, audio and cross-modal

One API covers all three modalities and scores a video on its frames AND its soundtrack, surfacing disagreement between them rather than averaging it away.

DURABILITY

Modular against the arms race

Detection is a moving target. Models are swappable behind a stable scoring interface, so a new generator is a model update, not a rewrite.



Upload. Score. Act.

Live product: React dashboard, documented REST API, ten models loaded, running today on image, video and audio.

1

Upload

Drag in a file, or POST to /api/v1. JPG PNG WebP, MP4 MOV MKV, WAV MP3 FLAC.

2

Analyse

Ten models score in parallel; learned fusion reconciles them with calibrated confidence.

3

Explain

Trust Score 0-100, GradCAM heatmap, per-model scores and an EXIF forensics report.

4

Act

Four-tier verdict, exportable report, and a one-click cybercrime complaint filing in India.

LIKELY MANIPULATED

POSSIBLY MANIPULATED

UNCERTAIN

LIKELY AUTHENTIC

Four-tier verdict — we say 'uncertain' when we are, rather than forcing a coin-flip.



Concept render of the analysis view. Replace with a live screenshot before pitching — the product is running.



What is measured, and what is not.

Every number below is reproducible from this repository. We label targets as targets.

Capability	Status	Measured result	Evidence
CorefakeNet unified image model	Shipped	82.5% acc · 90.9% ROC-AUC · F1 81.3	332 held-out samples
Single-image latency (CPU)	Shipped	564 ms — 4.9x faster than the ensemble	Benchmark run
Video pipeline, end to end	Shipped	Verified at 1/2/10/20 FPS, all 200 OK	docs section 18
Audio deepfake detection	Shipped	Wav2Vec2 backbone — 97.9% on the model card	Upstream benchmark
REST API + React dashboard	Live	10 models loaded, OpenAPI documented, rate-limited	Deployed build
7-model ensemble accuracy	In progress	95%+ target — fusion calibration being corrected	Not yet claimable
Paying customers	Not yet	Zero. Seeking design partners.	Honest baseline

We are pre-revenue. The ask funds the move from a working, measured prototype to a calibrated product with design partners under contract.

A \$170M market compounding at 47.6%.

Early, mandated, and accelerating fast.

\$5.6B

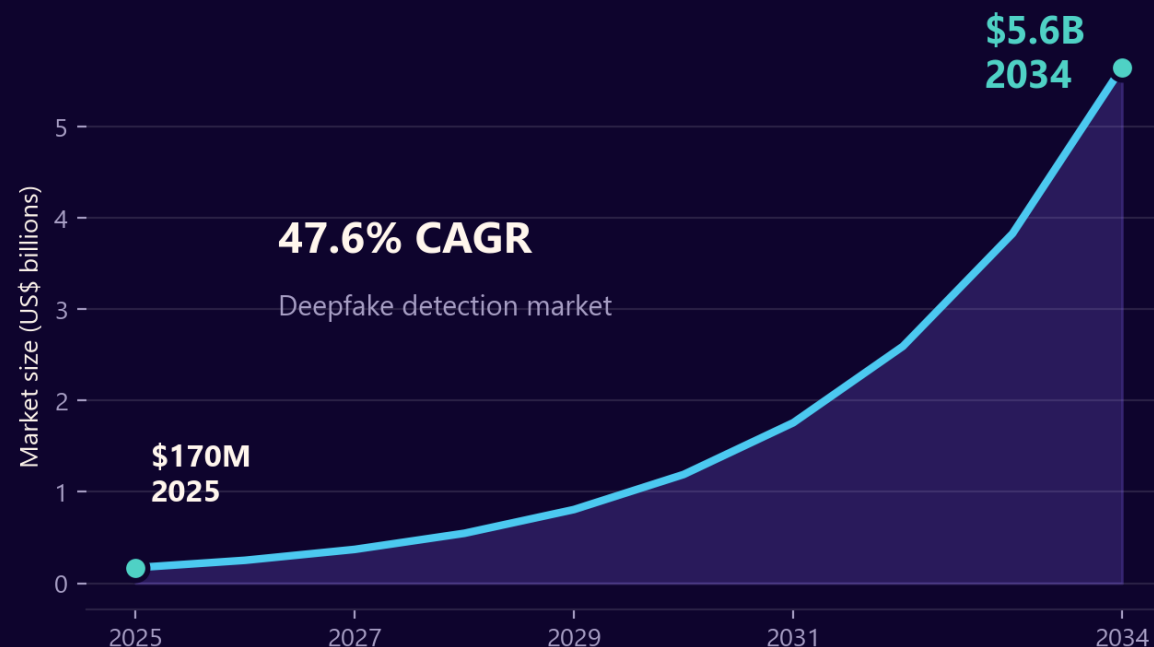
Deepfake detection market by 2034
(from \$170M in 2025)

66.7%

Of the market is video and image
detection — our core strength

49.2%

Media and journalism share, with
financial services fastest-growing



Adjacent budget lines we can reach from the same API: digital identity verification (~\$18B by 2027) and content moderation (~\$12B by 2027).



TAM, SAM and the share we are underwriting.

Sized bottom-up from the segments a detection API can actually invoice, not from the whole cybersecurity market.

WHERE THE FIRST DOLLARS COME FROM

TAM

Every org handling digital media, identity or communications (by 2035)

\$15-20B

Financial services

KYC/AML and voice-auth fraud. \$0.50-1.50 per verification. Fastest-growing vertical.

SAM

Organisations actively buying deepfake detection (by 2034)

\$5-8B

Media and newsrooms

Pre-publication verification. Largest current share of the market at ~49%.

SOM

ProofyX ARR at year 5 = 2-4% of SAM

\$100-200M

Government and legal

Evidence authentication and counter-disinformation. \$0.82B defence segment in 2025.

Comparable: Doppel grew enterprise customers 400% in one year.

A funded field — and a gap in it.

26 funded companies have raised \$259M between them. None of the leaders ship breadth and explainability in the same product.

Player	Raised	Modalities	Explainable	Fast mode
Reality Defender	\$52.4M	Multi	Partial	No
Doppel	\$124M	Social eng.	No	Real-time
Pindrop	Major VC	Voice only	No	Yes
GetReal Security	\$17.5M	Visual	Forensic	No
Sensity AI	Undisclosed	Visual	Forensic	No
Intel FakeCatcher	Internal	Video only	Physiological	Yes
ProofyX	Pre-seed	Image+Video+Audio	GradCAM + per-model	Yes (4.9x)

Our narrow, defensible claim:

ProofyX is the only stack combining all three modalities, per-model explainability and a fast single-pass mode behind one API — the shape a regulated buyer needs to file a decision, not just make one. We are not claiming the best single-model accuracy; we are claiming the most auditable answer.



Four ways to charge for the same API call.

Land through self-serve and the API, expand into enterprise licences. Benchmarked against what this market already pays.

Stream	Price	Target customer	Why they buy
API-as-a-service	\$0.25-0.50 / analysis	Platforms, integrators	Volume moderation, drop-in call
SaaS subscription	\$99-499 / month	Newsrooms, SMBs	Desk-level verification workflow
Per-check KYC	\$0.50-1.50 / check	Financial services	Regulatory duty, priced per user
Enterprise licence	\$50K-500K / year	Banks, government	On-prem, SLA, audit trail

BENCHMARK

Sightengine charges \$99/mo for 40,000 images;
deepidv \$19/mo + \$0.50/check.

GROSS MARGIN

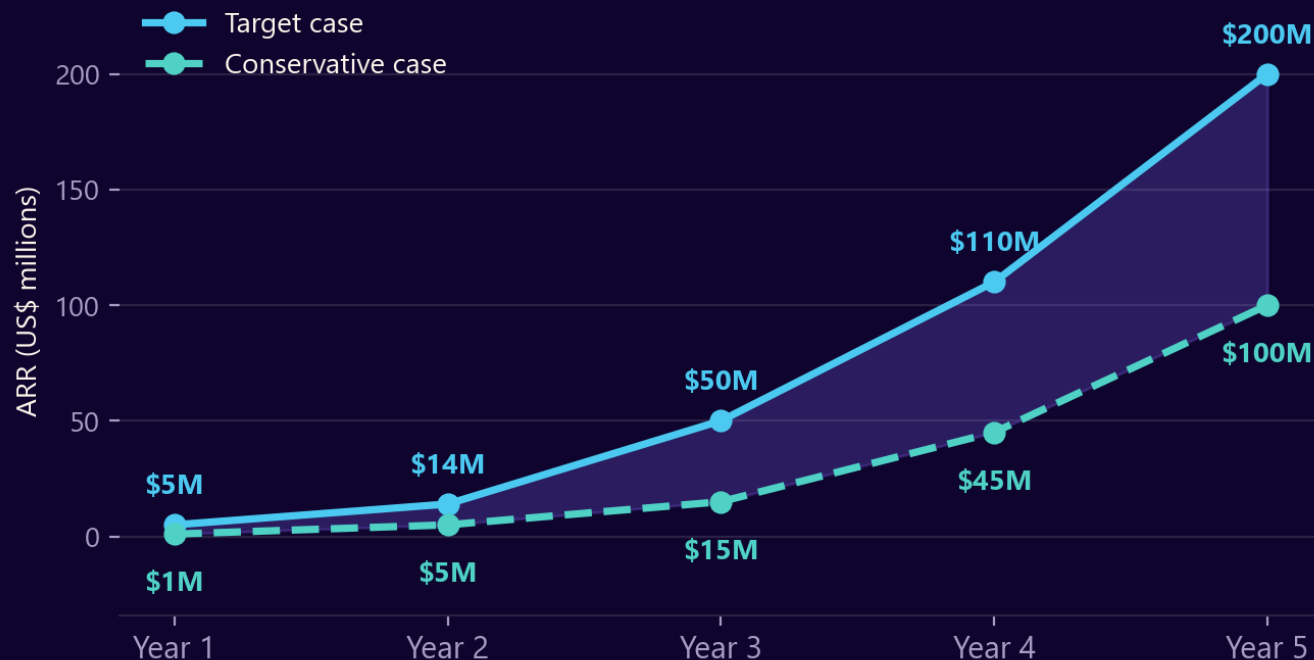
CPU inference today; GPU batch on the roadmap
moves unit cost down as volume rises.

EXPANSION

Same API serves moderation, KYC and evidence
— one integration, three budget lines.

Five-year path to \$100-200M ARR.

Conservative and target cases. Pre-revenue today — these are projections, and the assumptions are listed opposite.



ASSUMPTIONS BEHIND THE CURVE

Year 1

3-6 design partners at \$50-150K. No self-serve revenue assumed.

Year 2-3

Self-serve API opens; enterprise ACV rises as the EU AI Act deadline bites.

Year 4-5

2-4% of a \$5-8B SAM. Requires GPU inference and a compliance-certified deployment.

Cost base

Team and compute dominate; inference cost per call falls with batch GPU.

Builders who already shipped the hard part.

Ten models, three modalities and a live API — built by this team before raising a rupee.



Muhammed Sayeedur Rahman S

Founder / ML & Platform

Owns the detection stack and API: model training, fusion calibration, FastAPI backend and the compliance/traceability layer.



Mohamed Abrar

Co-founder / Product & Frontend

Largest contributor by commit count. Owns the React dashboard, the analysis workflows and the deployment pipeline.



Mohamed Shahid

Co-founder / Infrastructure & QA

Owns CI/CD, containerisation, security scanning and the test suite across the backend.

What we are missing, and are hiring for

A commercial lead with enterprise security or regtech sales experience, and a senior advisor from media forensics or financial-crime compliance. Both are line items in the ask.



Raising \$1.5M-\$2.5M seed.

18-24 months of runway to reach calibrated accuracy, paying design partners and an EU AI Act-ready deployment.

40%

Engineering

Two ML engineers and one platform engineer. Fix fusion calibration, ship GPU inference.

25%

Go-to-market

Commercial lead plus design-partner pilots in financial services and media.

20%

Compute & data

GPU training and inference, licensed evaluation datasets, red-team generation.

15%

Compliance

Security certification, audit trail and the evidence pack regulated buyers require.

WHAT THE MONEY BUYS — MILESTONES WE WILL BE JUDGED ON

Month 6

Ensemble calibration fixed; published accuracy on a public benchmark, third-party reproducible.

Month 12

3-6 design partners live under paid pilots; first enterprise contract signed.

Month 18

EU AI Act-ready deployment, SOC 2 in progress, GPU inference cutting unit cost.

Instrument: priced seed or SAFE. Valuation open — we would rather find the right partner than the highest cap.



Let's make digital evidence provable again.

We are looking for a lead investor and three design partners in financial services or media. A live demo takes ten minutes.


EMAIL REPLACE-WITH-REAL@EMAIL

REPOSITORY github.com/muhammedsayeedurrahman/ai-authenticity-checker

PHONE REPLACE-WITH-REAL-NUMBER

Investor Pitch Deck

August 8, 2026





Every number **has a source.**

Full research file: docs/PROOFYX_COMPLETE_ANALYSIS.md in the repository.

THREAT STATISTICS

- 8M+ deepfake files — Bright Defense
- \$3B+ US losses, 2025 — Security Magazine
- \$280K average incident — DeepStrike
- 0.1% human accuracy — iProov

MARKET SIZING

- \$170M to \$5.6B, 47.6% CAGR — Market.us
- Deepfake AI market — Mordor Intelligence
- \$15.1B by 2035 — OpenPR
- Defence \$0.82B — MarketsandMarkets

REGULATION

- EU AI Act Art. 50 / 99 — Blackbird.AI
- TAKE IT DOWN Act — Stack Cyber
- DoD investment — MIT Tech Review
- India IT Rules 2026 — docs/COMPLIANCE.md

COMPETITION AND FUNDING

- \$259M across 26 companies — Tracxn
- Doppel \$70M Series C — Fortune
- Reality Defender \$33M — SecurityWeek
- GetReal \$17.5M — TechCrunch

ProofyX measured results are reproducible from this repository: scripts/make_pitch_figures.py regenerates every chart in this deck from the recorded evaluation runs.